

# High-Throughput Request Gateway with Circuit Breaker

A lightweight, efficient API Gateway built using FastAPI that routes requests to backend microservices with round-robin load balancing and a built-in circuit breaker mechanism for fault tolerance.

## Features:

- Round-Robin Load Balancing across multiple service instances
- Per-URL Circuit Breaker to isolate failed services
- Configurable Service Mappings via YAML (`config.yaml`)
- Hot Configuration Reloading without restarting the gateway
- Cooldown & Retry Logic with half-open state recovery
- Built with FastAPI and httpx for performance and scalability

## Architecture Overview:

- The gateway sits in front of microservice replicas and handles:
- Request routing based on service name
- Load balancing using round-robin strategy
- Failure detection and circuit breaking
- Transparent recovery from failed services

Client —> Gateway (`/v1/proxy/{service}`) —> Healthy Backend  
└─X─> Skipped (Breaker Open)

## Project Structure:

File/Directory	Description
main.py	FastAPI app with routing and proxy logic
circuit_breaker.py	Circuit breaker implementation
config.yaml	Service-to-URL mapping configuration
requirements.txt	Python dependencies
dockerfile	Docker file commands
README.md	Project documentation

## Installation and Usage:

### 1. Install Dependencies:

```
pip install -r requirements.txt
```

### 2. Start Dummy Backend Services (Optional for Testing):

```
python -m http.server 9001
```

```
python -m http.server 9002
```

### 3. Run the Gateway:

```
uvicorn main:app --host 127.0.0.1 --port 8000
```

### 4. Send a Proxy Request:



```
curl http://127.0.0.1:8000/v1/proxy/user-service
```

### 5. Reload Config Without Restart:

```
curl -X POST http://127.0.0.1:8000/v1/reload-config
```

## Circuit Breaker Logic:

Each backend URL maintains its own circuit state:

- Closed: Normal traffic flow.
- Open: Skips traffic after 5 consecutive failures.
- Cooldown: Skips traffic for 60 seconds.
- Half-Open: Sends 1 test request.
  -  If success → circuit closes.
  -  If failure → circuit reopens.

This logic prevents request pileups to failed services and enables automatic recovery.

## Configuration File: config.yaml:

user-service:

- http://localhost:9001

- http://localhost:9002

**Testing Strategy:**

- Simulate failures by manually stopping one backend
- Validate load balancing with multiple replicas
- Observe circuit breaker triggering and recovery using logs or request responses